

- Stateless File Ingestion Gateway - Technical Report
  - System Overview
  - Architecture
    - Components
    - Key Files
  - Data Flow
    - State Transitions
    - Timing (Production vs Demo)
  - Recovery & Deduplication
    - Recovery Mechanism
    - Deduplication
  - Running the System
    - Start
    - Demo Script
    - Verify
  - Lessons Learned

# Stateless File Ingestion Gateway - Technical Report

---

## System Overview

---

A stateless file ingestion microservice that receives file uploads, stores them in MinIO, tracks state in Elasticsearch with automatic ILM-based cleanup, and publishes metadata to Kafka.

```
Parse error on line 1:  
flowchart LR      Cli  
^  
Expecting 'NEWLINE', 'SPACE', 'GRAPH', got 'ALPHA'
```

---

## Architecture

---

# Components

Component	Purpose	Port
ingestion-gateway	Spring Boot app handling uploads and background processing	8080
mongodb	Operational database for batch state	27017
minio	S3-compatible object storage for files	9000/9001
kafka	Message broker for downstream consumers	29092
mongo-express	MongoDB admin interface	8081
kafka-ui	Kafka message visualization	8090

## Key Files

File	Purpose
IngestionController.java	REST endpoint /api/batches/upload
IngestionService.java	Upload processing logic (PENDING → READY) with 5s delay
BatchProcessor.java	Background worker (READY → DONE + Kafka) with 5s delay
BatchDocument.java	MongoDB document model
MinioService.java	MinIO interactions

## Data Flow

## State Transitions



```
└── Files uploaded to MinIO  
└── Document created in MongoDB
```

## Timing (Production vs Demo)

Event	Production Timing	Demo Behavior
PENDING	Immediate	<b>5s Delay</b> (Simulated blocking)
READY	After upload	<b>5s Delay</b> (Simulated consumer latency)
DONE	Asynchronous	Immediate after processing
RECOVERY	N/A	Worker runs every <b>30s</b> , recovers batches > <b>5m</b> old

## Recovery & Deduplication

### Recovery Mechanism

- **Problem:** What if Kafka is down when **IngestionService** tries to publish the "fire-and-forget" event?
- **Solution:** The **BatchProcessor** runs a Scheduled Task every **30 seconds**.
- **Logic:** It queries MongoDB for batches that are in **READY** state for more than **5 minutes**. These are considered "stuck" and the event is re-published to Kafka.

### Deduplication

- **Strategy:** Content-based addressing in **tmp** bucket + UUIDs in **prod** bucket.
- **Tmp:** **tmp-bucket/yy/MM/dd/<sha256-hash>** (Deduplicates identical uploads).
- **Prod:** **prod-bucket/yy/MM/dd/<uuid>/<filename>** (Ensures unique delivery for downstream consumers).

## Running the System

# Start

```
docker-compose up -d --build
```

## Demo Script

```
.\demo-refactoring.ps1
```

## Verify

- **Kafka UI:** <http://localhost:8090> → Topics → `batch.ingestion.events`
  - **MinIO:** <http://localhost:9001> (minioadmin/minioadmin)
  - **Mongo Express:** <http://localhost:8081>
- 

## Lessons Learned

1. **Kafka Recovery:** Implementing a "pull" based recovery (scheduled task) complements "push" based events for robustness.
2. **Dual Bucket Strategy:** Decoupling ingestion (deduplicated tmp) from delivery (unique prod) saves storage while meeting delivery requirements.
3. **Demo Observability:** Adding artificial delays is crucial for demonstrating transient states (PENDING/READY) that would otherwise process too fast to see.